

Chat Filter Accuracy Audit

Ragnarok: Rebirth Global · independent analysis · 10 September 2026

The chat filter matches blocked entries as **raw substrings**. An entry therefore fires inside any longer word that happens to contain it, so a large volume of ordinary vocabulary is censored in normal conversation. The same design provides little resistance to deliberate evasion. Both effects follow from one implementation choice.

1. FINDING — ORDINARY WORDS ARE CENSORED AT SCALE

Measured against the title's own English localization, using the filter word list shipped in the client. Every string below was authored by the publisher, so each flag is unambiguously a false positive.

guild ← blocked by "ui"	skill ← "ll"	your ← "ur"	have ← "av"						
all	will	water	after	attack	monster	item	team	they	Prontera
Freezing → ****zing	Savage Card → **vage card								

Those are among the most frequent words in the title's own localization — **guild** appears 7,586 times, **skill** 6,109, **all** 3,484. **5,723 words that appear in the game's own text are censored by the game's own filter**, and 161,084 across general English vocabulary, place names and surnames. Each is supplied with the fragment responsible as evidence.

LOCALE	OWN ITEM & SKILL NAMES FLAGGED	
Portuguese	52.1%	Across all 81,303 localized English strings the rate is 69.0% ; against an ordinary English dictionary, 37.1% . The list is 95,660 entries, 93.4% Chinese — market-compliance material, appropriate for the market it was written for. 30,262 entries are one to three characters , and as substrings these sit inside ordinary vocabulary.
Indonesian	48.0%	
English	46.1%	
Thai	30.0%	
Chinese	8.4%	

2. FINDING — DETECTION OF DISGUISED ABUSE IS WEAK

Across ten common obfuscation strategies the list detects **22.9%**. Spaced-out text, full-width characters and homoglyph substitutions pass unmodified; of 25 common English profanity terms tested, **6** appear in the list at all. Correcting the matching improves accuracy in **both** directions — it is not a reduction in moderation coverage.

RECOMMENDATIONS

1. **Deploy the supplied allowlist.** A verified list of ordinary words that the current list censors — each one drawn from a dictionary, gazetteer or census file, with the responsible fragment recorded as evidence. Purely additive: it can only remove false positives, never add blocking. No matcher changes required.

2. **Apply a minimum fragment length.** One conditional: an entry shorter than N characters matches only as a whole word. Nothing is deleted and no entry is reviewed, and unlike a list it covers vocabulary that has never been enumerated — surnames, place names, future slang.

OWN NAMES FLAGGED	CURRENT	N=3	N=4
English	46.1%	7.6%	4.0%
Portuguese	52.1%	8.2%	0.9%
Indonesian	48.0%	6.2%	3.1%
Thai	30.0%	4.3%	2.4%

3. **Add the supplied CI check.** ~40 lines, no dependencies, enforcing one invariant: no string the publisher authored may be flagged by its own filter.

Scope and limitations. Analysis used the filter word list shipped in the client build, which does not exactly match the list enforced server-side and which changes over time. The conclusions about matching behaviour are independently supported by observed live behaviour, but specific entries named in the supporting material should be re-derived against the production list before action; all tooling to do so is included and every figure is reproducible. · No recommendation requires a client update — the filter check is already a server round trip.